

Projet 4_Analyse de ventes

Parcours Data Analyst

Xuefei ZHANG_2021

Readme

Ce projet est composé principalement de 3 parties: traitement et nettoyage, analyse univariée, et analyse bivariée.

Partie 1 se déroule comme suit:

- Import des librairies
- Import des CSV et création de 3 dataframes originaux: products, customers, transactions
- Création du dataframe principal: trans_prod_cus
- Nettoyage et traitement du dataframe principal
- Enrichissement du dataframe principal: création de nouvelles colonnes (âge, year, month, period...) pour faciliter l'exploitation et l'analyse

Partie 2 consiste à :

Réaliser des analyses univariées pour certaines variables du dataframe:

- Price
- Catégorie (categ)
- Age

Réaliser des analyses bivariées (qualitative et quantitative) entre les variables du dataframe:

- Age VS taille du panier moyen (nombre d'articles par session)
- Age VS Fréquence d'achat (nombre d'achats)
- Catégorie VS Sexe
- Age VS Montant d'achat
- Catégorie VS Montant d'achat
- Age VS Catégorie
- Age VS Montant d'achat

Sommaire

- ❑ Réorganisation des dataframes
et Nettoyage
- ❑ Analyse univariée
- ❑ Analyse bivariée

1. Dataframe principale et nettoyage

- Constitution de dataframe principale
- Nettoyage de données
- Traitement de données manquantes

Dataframes originaux :

customers , products , transactions

In [786]:

```
customers = pd.read_csv("customers.csv")  
  
print(customers.describe(include="all"))  
print(customers.isnull().values.any())  
customers.head()
```

	client_id	sex	birth
count	8623	8623	8623.000000
unique	8623	2	NaN
top	c_1565	f	NaN
freq	1	4491	NaN
mean	NaN	NaN	1978.280877
std	NaN	NaN	16.919535
min	NaN	NaN	1929.000000
25%	NaN	NaN	1966.000000
50%	NaN	NaN	1979.000000
75%	NaN	NaN	1992.000000
max	NaN	NaN	2004.000000
False			

Out[786]:

	client_id	sex	birth
0	c_4410	f	1967
1	c_7839	f	1975
2	c_1699	f	1984
3	c_5961	f	1962
4	c_5320	m	1943

In [787]:

```
products = pd.read_csv("products.csv")  
print(products.describe(include="all"))  
print(products.isnull().values.any())  
print(products.info())  
products.sort_values(by='price', ascending=True).head()
```

	id_prod	price	categ
count	3287	3287.000000	3287.000000
unique	3287	NaN	NaN
top	1_50	NaN	NaN
freq	1	NaN	NaN
mean	NaN	21.856641	0.370246
std	NaN	29.847908	0.615387
min	NaN	-1.000000	0.000000
25%	NaN	6.990000	0.000000
50%	NaN	13.060000	0.000000
75%	NaN	22.990000	1.000000
max	NaN	300.000000	2.000000
False			

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3287 entries, 0 to 3286
Data columns (total 3 columns):
Column Non-Null Count Dtype

0 id_prod 3287 non-null object
1 price 3287 non-null float64
2 categ 3287 non-null int64
dtypes: float64(1), int64(1), object(1)
memory usage: 77.2+ KB
None

Out[787]:

	id_prod	price	categ
731	T_0	-1.00	0
2355	0_202	0.62	0
2272	0_528	0.62	0
370	0_120	0.66	0
1211	0_1844	0.77	0

the product with id_prod = T_0 has its price negative, to clean

In [1096]:

```
transactions = pd.read_csv("transactions.csv")  
print(transactions.describe(include="all", sep='\n'))  
print(transactions.isnull().values.any())  
print(transactions.info())  
transactions.head()
```

	id_prod	date	session_id	client_id
count	337016	337016	337016	337016
unique	3266	336855	169195	8602
top	1_369	test_2021-03-01	02:30:02.237413	s_0 c_1609
freq	1081	13	200	12855
False				

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 337016 entries, 0 to 337015
Data columns (total 4 columns):
Column Non-Null Count Dtype

0 id_prod 337016 non-null object
1 date 337016 non-null object
2 session_id 337016 non-null object
3 client_id 337016 non-null object
dtypes: object(4)
memory usage: 10.3+ MB
None

Out[1096]:

	id_prod	date	session_id	client_id
0	0_1483	2021-04-10 18:37:28.723910	s_18746	c_4450
1	2_226	2022-02-03 01:55:53.276402	s_159142	c_277
2	1_374	2021-09-23 15:13:46.938559	s_94290	c_4270
3	0_2186	2021-10-17 03:27:18.783634	s_105936	c_4597
4	0_1351	2021-07-17 20:34:25.800563	s_63642	c_1242

Repérer et traiter les valeurs manquantes et aberrantes

- products

```
In [6]: products = products[products.id_prod != "T_0"]
print(products.info())
products.sort_values(by='price', ascending=True).head()

# the product with negative price is dropped from the dataframe
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 3286 entries, 0 to 3286
Data columns (total 3 columns):
#   Column  Non-Null Count  Dtype
---  -
0  id_prod  3286 non-null      object
1  price    3286 non-null      float64
2  categ    3286 non-null      int64
dtypes: float64(1), int64(1), object(1)
memory usage: 102.7+ KB
None
```

Le produit avec id_prod = T_0 est enlevé du dataframe

Out[6]:

	id_prod	price	categ
2272	0_528	0.62	0
2355	0_202	0.62	0
370	0_120	0.66	0
1211	0_1844	0.77	0
1530	0_1620	0.80	0

Repérer et traiter valeurs manquantes et aberrantes

- transactions

In [13]: *# Select all duplicate rows based on values of [date] column*

```
date_duplicate = transactions[transactions.duplicated(['date'])]

print(date_duplicate.describe(include="all"))

print("Duplicate rows based on the column 'date' are:")
date_duplicate.head()
```

	id_prod	date	session_id	client_id
count	161	161	161	161
unique	1	37	1	2
top	T_0	test_2021-03-01 02:30:02.237413	s_0	ct_0
freq	161	12	161	84

Duplicate rows based on the column 'date' are:

Out[13]:

	id_prod	date	session_id	client_id
27161	T_0	test_2021-03-01 02:30:02.237434	s_0	ct_0
34387	T_0	test_2021-03-01 02:30:02.237443	s_0	ct_0
48425	T_0	test_2021-03-01 02:30:02.237443	s_0	ct_1
54813	T_0	test_2021-03-01 02:30:02.237412	s_0	ct_1
56373	T_0	test_2021-03-01 02:30:02.237446	s_0	ct_0

it's clear that all these duplicated rows have some points in common: id_prod = T_0, date started by 'test_', session_id = 's_0'

In [16]: *#To drop duplicate rows in [date] column*

```
transac=transactions[transactions.id_prod != "T_0"]
print(transac.describe())
transac.head()
```

```
transac2 = transac[transac.session_id != "s_0"]
print(transac2.describe())
```

=> no difference in case we deduplicate by id_prod or by session_id

```
transac.head()
```

	id_prod	date	session_id	client_id
count	336816	336816	336816	336816
unique	3265	336816	169194	8600
top	1_369	2022-02-28 22:55:33.742567	s_118668	c_1609
freq	1081	1	14	12855

	id_prod	date	session_id	client_id
count	336816	336816	336816	336816
unique	3265	336816	169194	8600
top	1_369	2022-02-28 22:55:33.742567	s_118668	c_1609
freq	1081	1	14	12855

Out[16]:

	id_prod	date	session_id	client_id
0	0_1483	2021-04-10 18:37:28.723910	s_18746	c_4450
1	2_226	2022-02-03 01:55:53.276402	s_159142	c_277
2	1_374	2021-09-23 15:13:46.938559	s_94290	c_4270
3	0_2186	2021-10-17 03:27:18.783634	s_105936	c_4597
4	0_1351	2021-07-17 20:34:25.800563	s_63642	c_1242

Combiner 3 dataframes dans un seul: trans_prod_cus

In [19]: *# to merge 3 dataframes into trans_prod_cus via primary keys*

```
trans_prod_cus = pd.merge(trans_prod, trans_cus, on = ['date', 'id_prod', 'session_id', 'client_id'])  
print(trans_prod_cus.info())  
trans_prod_cus.head()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 336816 entries, 0 to 336815  
Data columns (total 8 columns):  
#   Column      Non-Null Count  Dtype  
--  --  
0   id_prod     336816 non-null  object  
1   date        336816 non-null  object  
2   session_id  336816 non-null  object  
3   client_id   336816 non-null  object  
4   price       336713 non-null  float64  
5   categ       336713 non-null  float64  
6   sex         336816 non-null  object  
7   birth       336816 non-null  int64  
dtypes: float64(2), int64(1), object(5)  
memory usage: 23.1+ MB  
None
```

Out[19]:

	id_prod	date	session_id	client_id	price	categ	sex	birth
0	0_1483	2021-04-10 18:37:28.723910	s_18746	c_4450	4.99	0.0	f	1977
1	2_226	2022-02-03 01:55:53.276402	s_159142	c_277	65.75	2.0	f	2000
2	1_374	2021-09-23 15:13:46.938559	s_94290	c_4270	10.71	1.0	f	1979
3	0_2186	2021-10-17 03:27:18.783634	s_105936	c_4597	4.20	0.0	m	1963
4	0_1351	2021-07-17 20:34:25.800563	s_63642	c_1242	8.99	0.0	f	1980

Conversion de types de data et création de nouvelles colonnes

In [21]: *#convert the date and birth columns from object into datetime*

```
trans_prod_cus['date'] = pd.to_datetime(trans_prod_cus.date)
trans_prod_cus['birth'] = pd.to_datetime(trans_prod_cus['birth'], format = '%Y', errors='coerce')

print(trans_prod_cus.info())
trans_prod_cus.head()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 336816 entries, 0 to 336815
Data columns (total 8 columns):
#   Column      Non-Null Count  Dtype
---  -
0   id_prod     336816 non-null object
1   date        336816 non-null datetime64[ns]
2   session_id  336816 non-null object
3   client_id   336816 non-null object
4   price       336713 non-null float64
5   categ       336713 non-null float64
6   sex         336816 non-null object
7   birth       336816 non-null datetime64[ns]
dtypes: datetime64[ns](2), float64(2), object(4)
memory usage: 23.1+ MB
None
```

Out[21]:

	id_prod	date	session_id	client_id	price	categ	sex
0	0_1483	2021-04-10 18:37:28.723910	s_18746	c_4450	4.99	0.0	f 197
1	2_226	2022-02-03 01:55:53.276402	s_159142	c_277	65.75	2.0	f 200
2	1_374	2021-09-23 15:13:46.938559	s_94290	c_4270	10.71	1.0	f 197
3	0_2186	2021-10-17 03:27:18.783634	s_105936	c_4597	4.20	0.0	m 196
4	0_1351	2021-07-17 20:34:25.800563	s_63642	c_1242	8.99	0.0	f 198

```
In [26]: trans_prod_cus['year-month-day'] = pd.to_datetime(trans_prod_cus['date']).dt.date
trans_prod_cus['year'] = pd.to_datetime(trans_prod_cus['date']).dt.year
trans_prod_cus['month'] = pd.to_datetime(trans_prod_cus['date']).dt.month
trans_prod_cus['period'] = trans_prod_cus['year'].astype(str) + trans_prod_cus['month'].astype(str)
trans_prod_cus['period'] = pd.to_numeric(trans_prod_cus['period'])

print(trans_prod_cus.info())
print(trans_prod_cus.isnull().values.any())
trans_prod_cus.head()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 336816 entries, 0 to 336815
Data columns (total 13 columns):
#   Column      Non-Null Count  Dtype
---  -
0   id_prod     336816 non-null object
1   date        336816 non-null datetime64[ns]
2   session_id  336816 non-null object
3   client_id   336816 non-null object
4   price       336713 non-null float64
5   categ       336713 non-null float64
6   sex         336816 non-null object
7   birth       336816 non-null datetime64[ns]
8   age         336816 non-null int64
9   year-month-day 336816 non-null object
10  year        336816 non-null int64
11  month       336816 non-null int64
12  period      336816 non-null int64
dtypes: datetime64[ns](2), float64(2), int64(4), object(5)
memory usage: 36.0+ MB
None
True
```

Out[26]:

	id_prod	date	session_id	client_id	price	categ	sex	birth	age	year-month-day	year	month	period
0	0_1483	2021-04-10 18:37:28.723910	s_18746	c_4450	4.99	0.0	f	1977-01-01	44	2021-04-10	2021	4	20214
1	2_226	2022-02-03 01:55:53.276402	s_159142	c_277	65.75	2.0	f	2000-01-01	21	2022-02-03	2022	2	20222
2	1_374	2021-09-23 15:13:46.938559	s_94290	c_4270	10.71	1.0	f	1979-01-01	42	2021-09-23	2021	9	20219
3	0_2186	2021-10-17 03:27:18.783634	s_105936	c_4597	4.20	0.0	m	1963-01-01	58	2021-10-17	2021	10	202110
4	0_1351	2021-07-17 20:34:25.800563	s_63642	c_1242	8.99	0.0	f	1980-01-01	41	2021-07-17	2021	7	20217

Repérer et traiter les valeurs null

```
In [28]: pricenull = pd.isnull(trans_prod_cus['price'])
categornull = pd.isnull(trans_prod_cus['categ'])
print(trans_prod_cus[pricenull].shape)
print(trans_prod_cus[categornull].shape)
```

Repérer des null dans les colonnes price et categ

```
(103, 13)
(103, 13)
```

=> then we find that the product with price and categ = NaN is product 0_2245.

given the id_prod starts always by the product categ, we can identify that product 0_2245 has categ 0.

```
In [29]: trans_prod_cus[pricenull]
```

Out[29]:

	id_prod	date	session_id	client_id	price	categ	sex	birth	age	year-month-day	year	month	period
6231	0_2245	2021-06-17 03:03:12.668129	s_49705	c_1533	NaN	NaN	m	1972-01-01	49	2021-06-17	2021	6	20216
10797	0_2245	2021-06-16 05:53:01.627491	s_49323	c_7954	NaN	NaN	m	1973-01-01	48	2021-06-16	2021	6	20216
14045	0_2245	2021-11-24 17:35:59.911427	s_124474	c_5120	NaN	NaN	f	1975-01-01	46	2021-11-24	2021	11	202111
17480	0_2245	2022-02-28 18:08:49.875709	s_172304	c_4964	NaN	NaN	f	1982-01-01	39	2022-02-28	2022	2	20222
21071	0_2245	2021-03-01 00:09:29.301897	s_3	c_580	NaN	NaN	m	1988-01-01	33	2021-03-01	2021	3	20213
21754	0_2245	2022-02-11 09:05:43.952857	s_163405	c_1098	NaN	NaN	m	1986-01-01	35	2022-02-11	2022	2	20222
22672	0_2245	2021-10-19 00:28:01.920054	s_106841	c_3953	NaN	NaN	f	1984-01-01	37	2021-10-19	2021	10	202110
24576	0_2245	2022-02-25 00:08:08.736068	s_170426	c_6236	NaN	NaN	f	1976-01-01	45	2022-02-25	2022	2	20222
30874	0_2245	2021-08-22 08:51:27.564509	s_79102	c_6752	NaN	NaN	m	1987-01-01	34	2021-08-22	2021	8	20218
31330	0_2245	2021-05-12 03:36:34.586221	s_33316	c_6205	NaN	NaN	f	1984-01-01	37	2021-05-12	2021	5	20215
34893	0_2245	2021-10-15 09:31:31.539354	s_105069	c_4188	NaN	NaN	f	1935-01-01	86	2021-10-15	2021	10	202110

À remarquer que pour tous ces valeurs null de colonne categ et price, id_prod est le même : 0_2245

=> Il faut trouver la categ et le price de ce produit dans dataframe products, et puis **imputation**

Traitement des valeurs manquantes dans colonnes price et categ

- Imputation

```
In [31]: # price in array of products categ 0
price0 = products.loc[products['categ']== 0 , 'price' ].values
price0
```

```
Out[31]: array([19.99,  5.13, 17.99, ..., 17.14, 11.22, 25.16])
```

```
In [33]: import statistics

print("Average price for products in categ 0:", round(statistics.mean(price0),2))
print("Median price for products in categ 0:", round(statistics.median(price0),2))
print("Mode of price for products in categ 0:", statistics.mode(price0))
```

Average price for products in categ 0: 11.73

Median price for products in categ 0: 10.32

Mode of price for products in categ 0: 4.99

```
In [25]: # to fill NaN in trans_prod_cus

trans_prod_cus.loc[trans_prod_cus['id_prod'] == "0_2245", ['price']] = trans_prod_cus['price'].fillna(10.32)
trans_prod_cus.loc[trans_prod_cus['id_prod'] == "0_2245", ['categ']] = trans_prod_cus['categ'].fillna(0)

print("after fillna, the number of null values in dataframe is:", trans_prod_cus.isnull().sum().sum())
trans_prod_cus.head()
```


after fillna, the number of null values in dataframe is: 0

Out[25]:

	id_prod	date	session_id	client_id	price	categ	sex	birth	age	year-month-day	year	month	period
0	0_1483	2021-04-10 18:37:28.723910	s_18746	c_4450	4.99	0.0	f	1977-01-01	44	2021-04-10	2021	4	20214
1	2_226	2022-02-03 01:55:53.276402	s_159142	c_277	65.75	2.0	f	2000-01-01	21	2022-02-03	2022	2	20222
2	1_374	2021-09-23 15:13:46.936559	s_94290	c_4270	10.71	1.0	f	1979-01-01	42	2021-09-23	2021	9	20219
3	0_2186	2021-10-17 03:27:18.783634	s_105936	c_4597	4.20	0.0	m	1963-01-01	58	2021-10-17	2021	10	202110
4	0_1351	2021-07-17 20:34:25.800563	s_63642	c_1242	8.99	0.0	f	1980-01-01	41	2021-07-17	2021	7	20217

In dataframe 'products', there is NOT a product with id_prod = 0_2245

- but it's sure that it is in categ 0
- thus for price, need to decide among mean/mode/median, which to use in categ 0



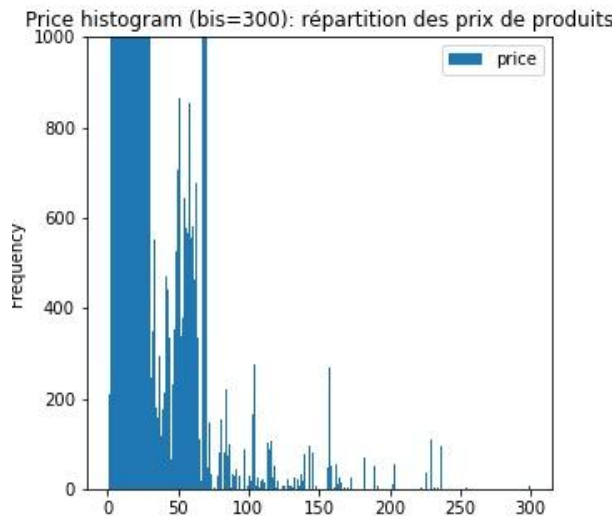
Based on this histogram, it's appropriate to use **median** of price in categ0

Analyse univariée

- Price
- Catégorie (categ)
- Age


Univariée - price - histogramme et kde

```
In [119]: trans_prod_cus.plot(kind="hist",  
    y="price",  
    figsize=(5,5),  
    bins=300,  
    ylim=(0,1000))  
  
plt.savefig('uni_categ_hist.jpg')
```



Répartition des prix de tous produits avec Kde

```
! sns.set(style="darkgrid")  
sns.histplot(data=trans_prod_cus, x="price", bins=300, kde = True)  
plt.show()  
plt.savefig('uni_price_hist_kde.jpg')
```



Mesure d'asymétrie

Left-skewed: price se concentre fortement dans l'éventail 0 - 30

Univariée - price - courbe Lorenz et indice Gini

In [136]:

```
depenses2 = trans_prod_cus['price'].values
n2 = len(depenses2)
lorenz2 = np.cumsum(np.sort(depenses2)) / depenses2.sum()
lorenz2 = np.append([0], lorenz2)  # La courbe de Lorenz commence à 0

fig, ax = plt.subplots(figsize=[5,5])
plt.axes().axis('equal')
xaxis = np.linspace(0-1/n, 1+1/n, n+1)
plt.title("Lorenz - concentration des prix")
plt.plot(xaxis, lorenz2, drawstyle='steps-post', color = "darkblue")
plt.plot(xaxis, xaxis, color = "blue")


plt.savefig("lorenz_courbe_price2.jpg")
```

In [138]:

```
AUC2 = (lorenz2.sum() -lorenz2[-1]/2 -lorenz2[0]/2)/n2
S2 = 0.5 - AUC2
gini2 = 2*S2
round(gini2,4)
```

Out[138]: 0.3921

Lorenz - concentration des prix



indice de Gini = 39.21%:

40% du montant de chiffres d'affaires viennent de 60% (100%-40%) des clients.

Selon la loi de Pareto (20/80), qui est la normalité, ici 40/60 est moins concentré que 20/80

=> une concentration pas très forte et proche égalitaire.

Univariée - catégorie (categ) - pie chart et boxplot

```
j: uni_cat=trans_prod_cus[["categ",'price']].groupby(by=["categ"])
uni_cat.head()
sumunicateg=uni_cat.sum()
display("mean of price by categ", uni_cat.mean())
display("sum of price by categ", uni_cat.sum())
display("median by categ:", uni_cat.median())
```

```
'mean of price by categ'
```

price	
categ	
0.0	10.646668
1.0	20.480106
2.0	75.174949


'sum of price by categ'

	price
categ	
0.0	2.230786e+06
1.0	2.247384e+06
2.0	1.319471e+06

```
'median by categ:'
```


price	
categ	
0.0	9.99
1.0	19.08
2.0	62.83

```
] plot = sumunicateg.plot.pie(subplots=True, figsize=(10, 6))
plt.ylabel("Montant total des ventes par catégorie")
plt.savefig('sum_uni_categ_pie.jpg')
```



- **Categ 2 apporte relativement moins de CA par rapport à categ 0 et 1;**
- **Prix de produits (moyenne et mediane) pour categ 0 augmentent en fonction de catégorie (0 - 1 - 2)**

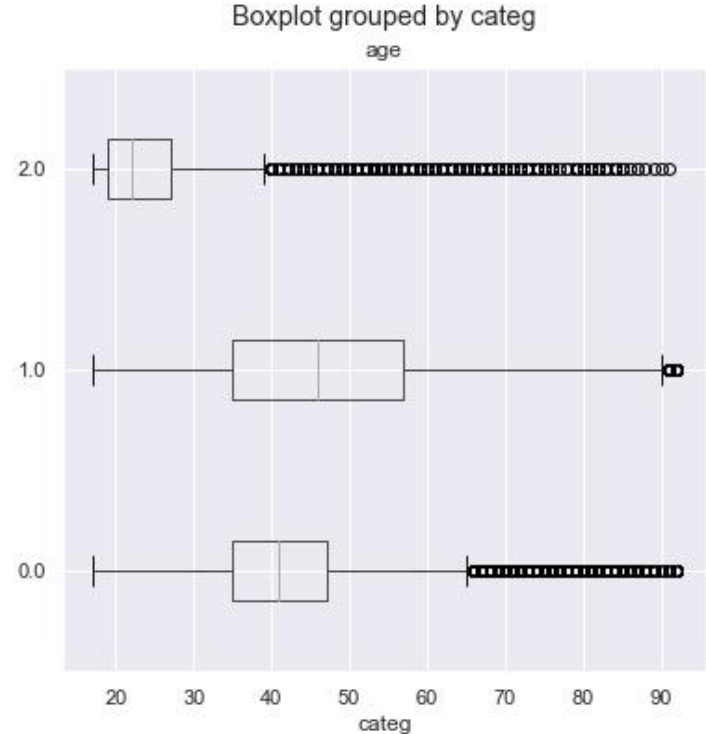
Boxplot grouped by categ



Univariée - âge vs categ


```
trans_prod_cus.boxplot(column = 'age', by = 'categ', vert=False, figsize=(6,6))
```

Outliers des âges ne sont pas des valeurs mauvaises, vu qu'en réalité les produits sont accessibles par la population de tout âge.



Univariée - âge - histogramme

```
sns.set(style="darkgrid")
sns.histplot(data=trans_prod_cus, x="age", bins=300, kde = True)
plt.show()
plt.savefig('uni_age_hist_kde.jpg')
```



Densité de l'âge:

La plupart des clients sont de tranche d'âge 30 - 55

Analyse bivariée


- Age VS taille du panier moyen (nombre d'articles par session)
- Age VS Fréquence d'achat (nombre d'achats)
- Catégorie VS Sexe
- Age VS Montant d'achat
- Catégorie VS Montant d'achat
- Age VS Catégorie
- Age VS Montant d'achat

Evolution CA au fil du temps _série temporelle

```
periode=pd.pivot_table(trans_prod_cus, index=['year', 'month'], columns=['categ'], values=['price'], aggfunc='sum')
periode.head()
```

Out[34]:

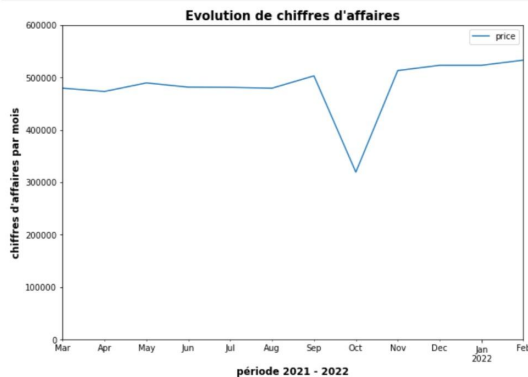
		price			
		categ	0.0	1.0	2.0
year	month				
2021	3		193659.26	186974.17	98771.48
	4		205304.15	156138.35	111682.70
	5		196197.52	165893.40	127359.59
	6		167958.58	189162.04	124209.56
	7		144753.20	188523.27	147663.47



```
In [37]: plt.figure();
periode.plot.bar(figsize=(10,7), ylim=(0, 600000), stacked=True);
plt.xlabel('period');
plt.ylabel('categ');

colors = {'category 0':'C0', 'category 1':'C1', 'category 2':'C2'}
labels = list(colors.keys())
handles = [plt.Rectangle((0,0),1,1, color=colors[label]) for label in labels]
plt.legend(handles, labels)
plt.title("Evolution des chiffres d'affaires par catégorie et par période", fontsize=12, fontweight='bold')
plt.show()
plt.savefig('Evolution des CA.png')
```

<Figure size 432x288 with 0 Axes>




Evolution de séances au fil du temps_série temporelle

```
periode2=pd.pivot_table(trans_prod_cus, index=['year', 'month'], columns=['categ'], values=['session_id'], aggfunc='count')
periode2.head()
print(periode2.shape)
periode2.head()
```

(12, 3)

		session_id			
		categ	0.0	1.0	2.0
year	month				
2021	3	18119	9134	1315	
	4	19335	7579	1501	
	5	18485	8107	1653	
	6	15886	9264	1669	
	7	13569	9169	1978	





A remarquer une anomalie:

Les chiffres d'affaires et les séances d'achat pour les produits de **catégorie 1** en **octobre 2021** sont tous relativement moins par rapport aux mois précédents et suivants

Que s'est-il passé avec les données d'octobre 2021?

```
month_pvt = trans_prod_cus.pivot_table(columns = ['year','month'], index = ['categ'], values = "price", aggfunc="sum")
month_pvt.head()
```

year	2021											2022
month	3	4	5	6	7	8	9	10	11	12	1	
categ												
0.0	193659.26	205304.15	196197.52	167958.58	144753.20	167770.70	246388.05	199290.18	155946.98	206048.68	164214.	
1.0	186974.17	156138.35	165893.40	189162.04	188523.27	162991.38	190613.78	33762.32	252910.39	251026.75	256267.	
2.0	98771.48	111682.70	127359.59	124209.56	147663.47	148635.99	65893.29	86179.70	104136.00	65934.49	102524	



Données du mois octobre 2021


```
In [76]: month10_pvt = month10.pivot_table(index = ['year-month-day'], columns = ['categ'], values = "price", aggfunc='sum')
print(month10_pvt.sum())
print(month10_pvt.head())

month10_pvt.plot(kind="bar",
                 figsize=(18,5),
                 stacked = False)

plt.savefig("month10.jpg")
```

```
categ
0.0    199290.18
1.0     33762.32
2.0     86179.70
dtype: float64
```

categ	0.0	1.0	2.0
year-month-day			
2021-10-01	6947.51	7003.79	2958.06
2021-10-02	7138.02	NaN	1895.13
2021-10-03	6783.58	NaN	2060.49
2021-10-04	6551.25	NaN	2600.09
2021-10-05	6357.91	NaN	3032.55
2021-10-06	7543.59	NaN	1798.12
2021-10-07	6404.01	NaN	1787.07
2021-10-08	7069.53	NaN	3137.82
2021-10-09	6808.69	NaN	2616.67
2021-10-10	6487.99	NaN	2188.68
2021-10-11	7005.40	NaN	3225.16
2021-10-12	6703.98	NaN	2118.19



Données manquantes sur categ 1 pour la période de 02 à 27 oct 2021

Chercher la normalité: données de categ 1 de tous les mois


```
In [69]: categ1 = month_pvt2[1].values
print(categ1)

import statistics
print('')
print('Statistics for products of category 1: ')
print("mean:", statistics.mean(categ1))
print("median:", statistics.median(categ1))
print("mode:", statistics.mode(categ1))

# statistics (mean, median, mode) by month

[186974.16999999 156138.35    165893.4    189162.03999999
 188523.26999999 162991.38    190613.77999999 33762.32
 252910.38999998 251026.74999998 256267.91999998 213120.63999999]

Statistics for products of category 1:
mean: 187282.03416665728
median: 188842.65499999246
mode: 186974.16999999323
```



tendance centrale: distribution unimodale et fortement centré autour de 188,000

Quelle valeur à imputer pour oct 2021?

```
In [31]: month_pvt2 = trans_prod_cus.pivot_table(columns = ['categ'], index = ['year', 'month'], values = "price", aggfunc = 'sum' )
month_pvt2
```

price amount subtotal of products in every categ / distribution de chiffres d'affaires par catégorie de produit -- bar chart

Out[31]:

		categ	0.0	1.0	2.0
year	month				
2021	3		193659.26	186974.17	98771.48
	4		205304.15	156138.35	111682.70
	5		196197.52	165893.40	127359.59
	6		167958.58	189162.04	124209.56
	7		144753.20	188523.27	147663.47
	8		167770.70	162991.38	148635.99
	9		246388.05	190613.78	65893.29
	10		199290.18	33762.32	86179.70
	11		155946.98	252910.39	104136.00
	12		206048.68	251026.75	65934.49
2022	1		164214.27	256267.92	102524.72
	2		183254.04	213120.64	136479.72

moyenne de prix de categ1 au courant de 12 mois	187282
total de prix de categ1(1 oct + 28-31 oct)	33762
moyenne de prix pour 02-27 oct	$=(187282 - 33762)/26 = 5905$

=> dans la table croisée de month10_pvt, il reste à imputer les manquantes de categ1 oct par **5905**


Imputation des données manquantes oct 2021

```
month10_pvt.reset_index(inplace = True)
month10_pvt[1].fillna(5905, inplace=True)
print(month10_pvt.sum())
month10_pvt.head()
```

```
categ
0.0  199290.18
1.0  187292.32
2.0   86179.70
dtype: float64
```

75]:

categ	0.0	1.0	2.0
year-month-day			
2021-10-01	6947.51	7003.79	2958.06
2021-10-02	7138.02	5905.00	1895.13
2021-10-03	6783.58	5905.00	2060.49
2021-10-04	6551.25	5905.00	2600.09
2021-10-05	6357.91	5905.00	3032.55



Data octobre après imputation

Analyse bivariée -

Âge vs nombre d'articles par séance (taille du panier moyen)

```
In [78]: age_panier = trans_prod_cus[['age', 'id_prod']].groupby(by=['age']).count()
age_panier.reset_index(inplace=True)
#age_panier.head()

# how many products brought by clients of same age
```

```
In [79]: customers['birth'] = pd.to_datetime(customers['birth'], format='%Y', errors='coerce')
customers['age'] = customers['birth'].apply(calculate_age)


age_pop = customers[['age', 'client_id']].groupby('age').count()
age_pop.reset_index(inplace=True)
#age_pop.head()

# pop of same age
```

```
In [80]: taille_panier = pd.merge(age_panier, age_pop, on='age')
taille_panier['panier_moyen'] = round(taille_panier['id_prod']/taille_panier['client_id'], 2)
taille_panier.head()
```

Out[80]:

	age	id_prod	client_id	panier_moyen
0	17	7348	440	16.70
1	18	2182	146	14.95
2	19	2224	146	15.23
3	20	2032	129	15.75
4	21	2175	136	15.99



Analyse bivariable -

Âge vs montant d'achats par séance

```
In [96]: age_session = trans_prod_cus[['age', 'session_id']].groupby(by=['age']).count()
age_session = age_session.rename(columns={'session_id': 'total_session'})
age_session.head()


age_price = trans_prod_cus[['age', 'price']].groupby(by=['age']).sum()
age_price = age_price.rename(columns={'price': 'total_montant'})
age_price.head()

age_montant = age_session.merge(age_price, how='left', on='age').merge(age_pop, on='age', how='left')

age_montant['montant_by_session'] = round(age_montant['total_montant']/age_montant['total_session'], 2)
age_montant['mean_session'] = round(age_montant['total_session']/age_montant['client_id'], 2)
age_montant.head()
```


Out[96]:

	age	total_session	total_montant	client_id	montant_by_session	mean_session
0	17	7348	295387.98	440	40.20	16.70
1	18	2182	88461.39	146	40.54	14.95
2	19	2224	89920.34	146	40.43	15.23
3	20	2032	84881.48	129	41.77	15.75
4	21	2175	80110.24	136	36.83	15.99



montant d'achats par séance
(montant by session)

Analyse bivariable - Âge vs fréquence d'achats



fréquence d'achats (mean session)

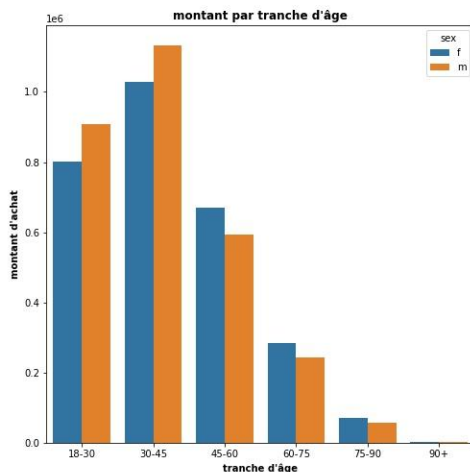
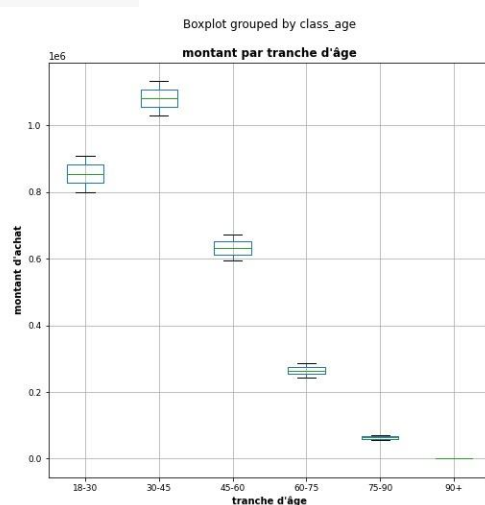
=> Clients de l'âge 30 à 53 connaissent des fréquences les plus élevées

Analyse bivariée - montant d'achat vs âge et sexe

```
def class_age(x):  
    if x <= 30:  
        return '18-30'  
  
    elif x > 30 and x <= 45:  
        return '30-45'  
  
    elif x > 45 and x <= 60:  
        return '45-60'  
  
    elif x > 60 and x <= 75:  
        return '60-75'  
  
    elif x > 75 and x <= 90:  
        return '75-90'  
  
    elif x > 90:  
        return '90+'  
  
trans_prod_cus['class_age'] = trans_prod_cus['age'].apply(lambda x: class_age(x))  
agesex_pvt = trans_prod_cus.pivot_table(index = ['class_age', 'sex'], values = "price",  
agesex_pvt.head()
```

Out[121]:

price		
class_age	sex	
18-30	f	800863.59
	m	908630.70
30-45	f	1029260.12
	m	1132994.96
45-60	f	671398.73




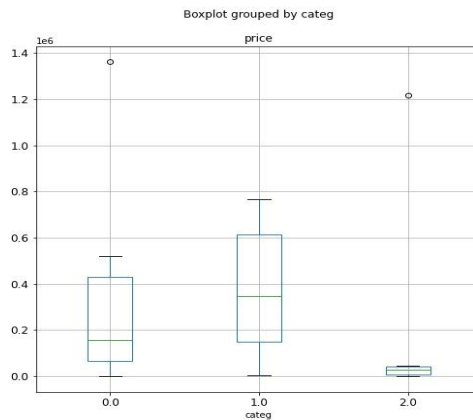
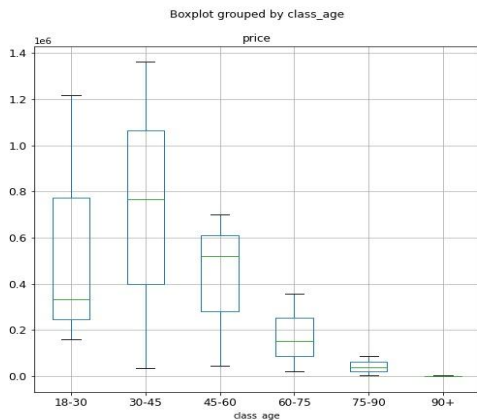
Analyse bivariée -

Âge vs montant, categ vs montant, Âge vs montant by categ

```
trans_prod_cus['class_age'] = trans_prod_cus['age'].apply(lambda x: class_age(x))  
  
agecateg_pvt = trans_prod_cus.pivot_table(index = ['class_age', 'categ'], values = 'price', aggfunc = sum)  
agecateg_pvt.head()
```

[13]:

price		
class_age	categ	
18-30	0.0	1.594341e+05
	1.0	3.342012e+05
	2.0	1.215859e+06
30-45	0.0	1.362260e+06
	1.0	
	2.0	




Analyse bivariable - sexe vs categ

```
In [100]: categ_sex = trans_prod_cus.pivot_table(index = ['categ'], columns = ['sex'], values = "price", aggfunc = sum)
          categ_sex
```

Out[100]:

sex	f	m
categ		
0.0	1104085.17	1126700.44
1.0	1137856.49	1109527.92
2.0	617799.31	701671.40

```
In [102]: plot = categ_sex.plot.pie(subplots=True, figsize=(10, 6))
          plt.savefig('categsex.jpg')
```




Par le montant d'achat

```
In [103]: categ_sex2 = trans_prod_cus.pivot_table(index = ['categ'], columns = ['sex'], values = "id_prod", aggfunc = count)
          categ_sex2
```

Out[103]:




sex	f	m
categ		
0.0	103846	105683
1.0	55469	54266
2.0	8260	9292

```
In [104]: plot = categ_sex2.plot.pie(subplots=True, figsize=(10, 6))
          plt.savefig('categsex2.jpg')
```



Par le nombre de produits achetés
(quantité d'achat)

Analyse bivariée - sex vs categ



Courbe Lorenz et indice Gini

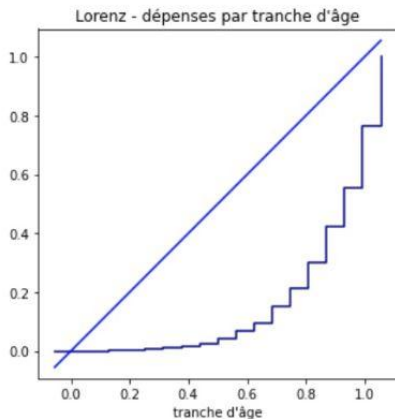
- concentration des dépenses en fonction de tranche d'âge

```
In [134]: depenses = agecateg_pvt['price'].values    #sous-totaux des dépenses par age (all people of same age) et par categ

n = len(depenses)
lorenz = np.cumsum(np.sort(depenses)) / depenses.sum()
lorenz = np.append([0], lorenz)    # La courbe de Lorenz commence à 0

fig, ax = plt.subplots(figsize=[5,5])
plt.axes().axis('equal')
xaxis = np.linspace(0-1/n, 1+1/n, n+1)    ## y a un segment de taille n pour chaque individu, plus 1 segment supplémentaire d'ordonnée 0. Le premier segment
plt.xlabel("tranche d'âge")
plt.title("Lorenz - dépenses par tranche d'âge")
plt.plot(xaxis, lorenz, drawstyle='steps-post', color = "darkblue" )
plt.plot(xaxis, xaxis, color = "blue")

plt.savefig("lorenz_courbe.jpg")
```



```
: AUC = (lorenz.sum() -lorenz[-1]/2 -lorenz[0]/2)/n    # Surface sous la courbe de Lorenz.
S = 0.5 - AUC    # surface entre la première bissectrice et le courbe de Lorenz
gini = 2*S
gini

: 0.6480633354879031
```

indice de Gini = 64%

=> au moins 64% du montant de chiffres d'affaires de clients de tranches d'âge 30-45 et 45-60 soit 36% de toute la population clientèle



Q & A



Merci

@Xuefei ZHANG 2021